

523454

# Computer Network Programming

Overview of TCP connections

Dr. Parin Sornlertlamvanich,

parin.s@sut.ac.th

# Socket Address Structures

## ■ IPv4 Socket Address Structure

```
struct in_addr {  
    in_addr_t    s_addr;        /* 32-bit IPv4 address */  
                                /* network byte ordered */  
};  
  
struct sockaddr_in {  
    uint8_t      sin_len;        /* length of structure (16) */  
    sa_family_t  sin_family;     /* AF_INET */  
    in_port_t    sin_port;       /* 16-bit TCP or UDP port number */  
                                /* network byte ordered */  
    struct in_addr sin_addr;     /* 32-bit IPv4 address */  
                                /* network byte ordered */  
    char         sin_zero[8];    /* unused */  
};
```

# Socket Address Structures

## ■ Generic Socket Address Structure

```
struct sockaddr {  
    uint8_t          sa_len;  
    sa_family_t      sa_family;    /* address family: AF_XXX value */  
    char             sa_data[14];  /* protocol-specific address */  
};
```

```
int bind(int, struct sockaddr *, socklen_t);  
struct sockaddr_in serv;          /* IPv4 socket address structure */
```

```
/* fill in serv */
```

```
bind(sockfd, (struct sockaddr *) &serv, sizeof(serv));
```

# Socket Address Structures

## ■ IPv6 Socket Address Structure

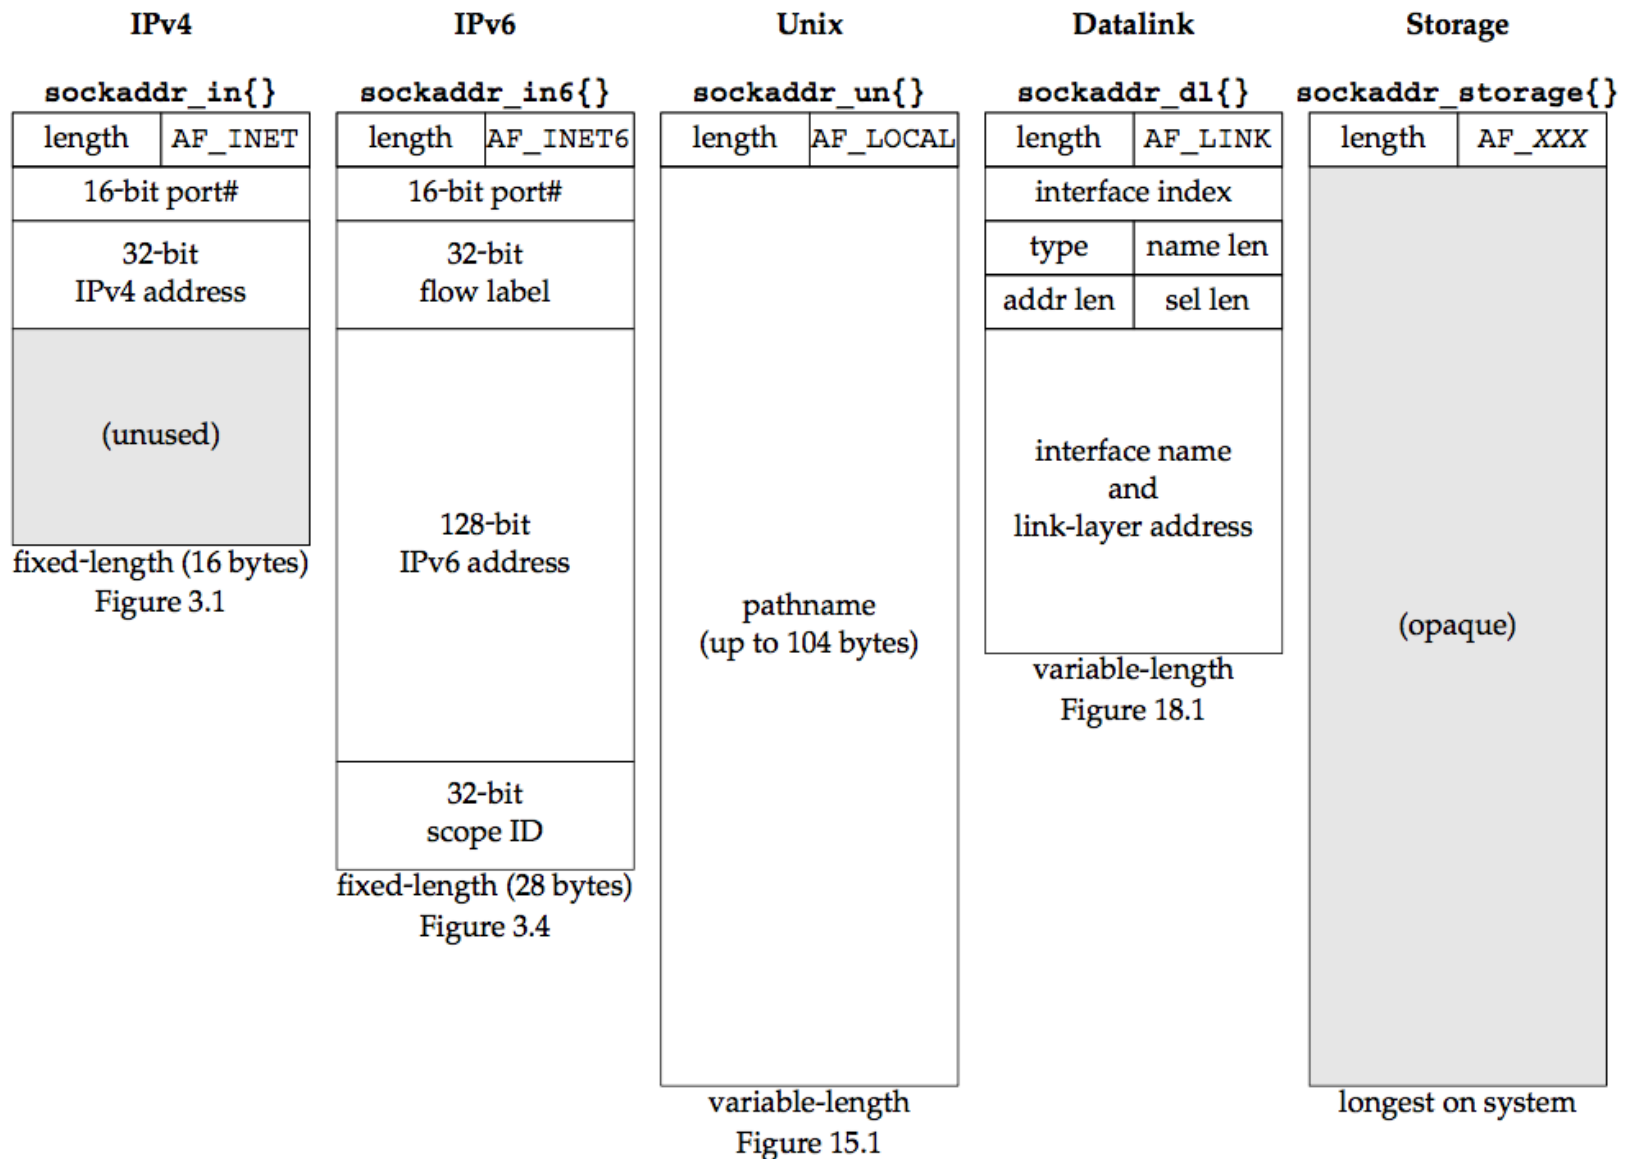
```
struct in6_addr {  
    uint8_t s6_addr[16];           /* 128-bit IPv6 address */  
                                    /* network byte ordered */  
};  
  
struct sockaddr_in6 {  
    uint8_t          sin6_len;      /* length of this struct (28) */  
    sa_family_t      sin6_family;   /* AF_INET6 */  
    in_port_t        sin6_port;     /* transport layer port# */  
                                    /* network byte ordered */  
    uint32_t          sin6_flowinfo; /* flow information, undefined */  
    struct in6_addr    sin6_addr;    /* IPv6 address */  
                                    /* network byte ordered */  
    uint32_t          sin6_scope_id; /* set of interfaces for a scope */  
};
```

# Socket Address Structures

## ■ New Generic Socket Address Structure

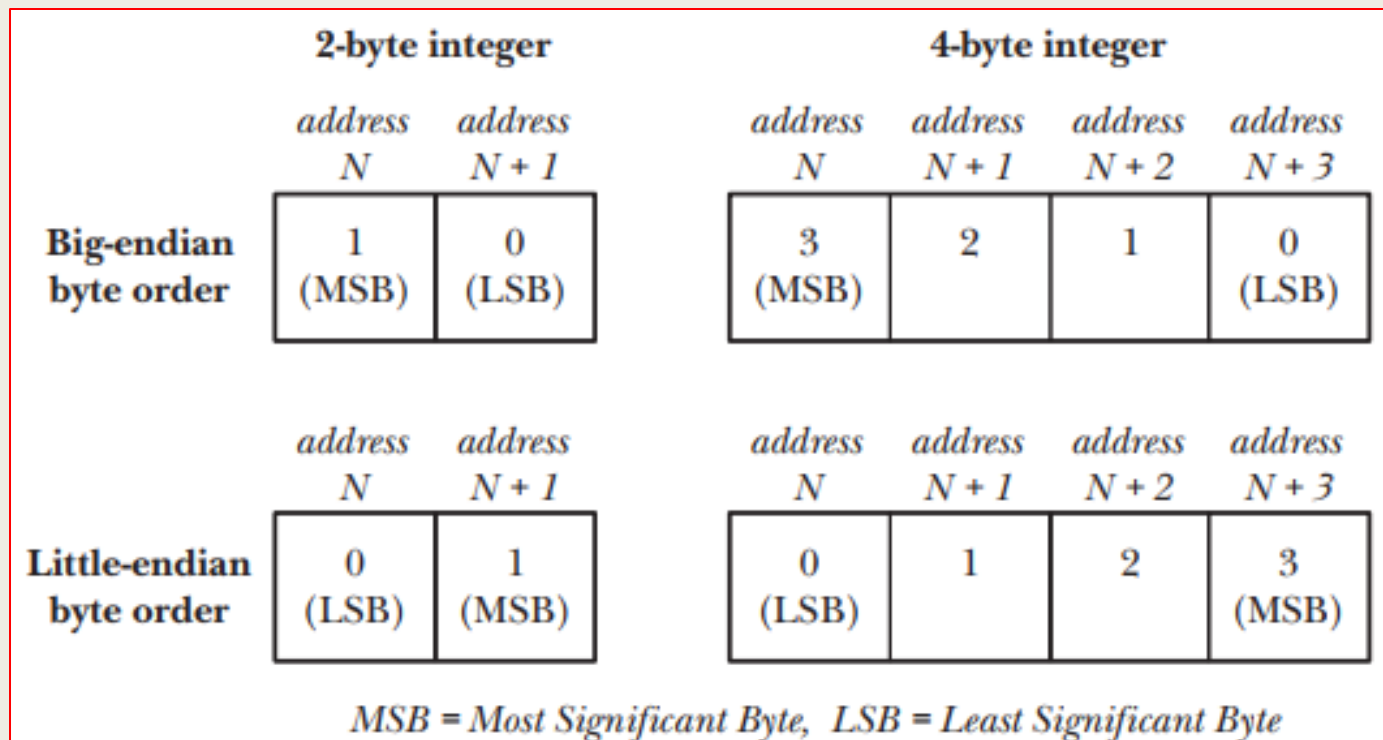
```
struct sockaddr_storage {  
    uint8_t          ss_len;           /* length of this struct (implementation  
                                       dependent) */  
    sa_family_t      ss_family;        /* address family: AF_xxx value */  
  
    /* implementation-dependent elements to provide:  
    * a) alignment sufficient to fulfill the alignment requirements of  
    *    all socket address types that the system supports.  
    * b) enough storage to hold any type of socket address that the  
    *    system supports.  
    */  
};
```

# Comparison of Socket Address Structures



# Byte Ordering Functions

- For a 16-bit integer that is made up of 2 bytes
  - there are two ways to store the two bytes in memory
- 1. Little-endian order: low-order byte is at the starting address.
- 2. Big-endian order: high-order byte is at the starting address.



# Byte Ordering Functions

	Low address				High address			
Address	0	1	2	3	4	5	6	7
Little-endian	Byte 0	Byte 1	Byte 2	Byte 3	Byte 4	Byte 5	Byte 6	Byte 7
Big-endian	Byte 7	Byte 6	Byte 5	Byte 4	Byte 3	Byte 2	Byte 1	Byte 0
Memory content	0x11	0x22	0x33	0x44	0x55	0x66	0x77	0x88
64 bit value on Little-endian				64 bit value on Big-endian				
0x8877665544332211				0x1122334455667788				


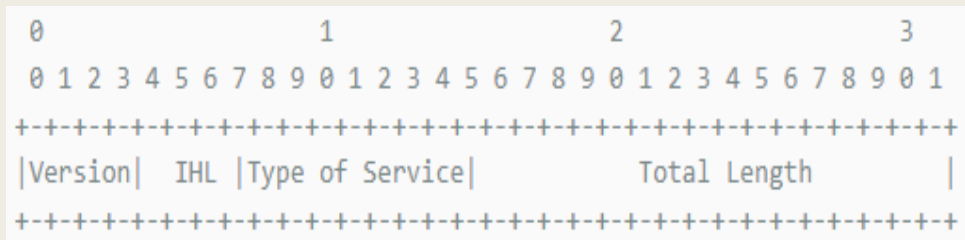


# Byte Ordering Functions

- The terms "little-endian" and "big-endian" indicate which end of the multibyte value, the little end or the big end, is stored at the starting address of the value.
- **Host byte order** refer to the byte ordering used by a given system.
  - Little-endian byte ordering
- Networking protocols must specify a **Network byte order**.
  - The Internet protocols use **big-endian byte ordering** for these multibyte integers

# Byte Ordering Functions

- Bit ordering is an important convention in Internet standards, such as the first 32 bits of the IPv4 header from RFC 791:



- the leftmost bit is the most significant. However, the numbering starts with zero assigned to the most significant bit.

# Byte Ordering Functions

```
#include <arpa/inet.h>
```

```
uint16_t htons(uint16_t host_uint16);
```

Returns *host\_uint16* converted to network byte order

```
uint32_t htonl(uint32_t host_uint32);
```

Returns *host\_uint32* converted to network byte order

```
uint16_t ntohs(uint16_t net_uint16);
```

Returns *net\_uint16* converted to host byte order

```
uint32_t ntohl(uint32_t net_uint32);
```

Returns *net\_uint32* converted to host byte order

- **h** stands for host
- **n** stands for network
- **s** stands for short (16-bit value, e.g. TCP or UDP port number)
- **l** stands for long (32-bit value, e.g. IPv4 address)

# Byte Ordering Functions

- To specify an endpoint address 158.182.9.1:5678

```
#include <winsock2.h>
...
struct sockaddr_in addr;
...
addr.sin_family = AF_INET;
addr.sin_port = htons(5678);
addr.sin_addr.s_addr = htonl(2662729985);
```

There are several ways to specify an IP address:

- Use a string – Good!
  - “158.182.9.1”
- Use an integer – Bad!
  - 2662729985
- Use DNS service – Good!

Remark:  $2662729985 = \begin{array}{cccc} 10011110 & 10110110 & 00000100 & 10000000 \\ \hline 158 & 182 & 9 & 1 \end{array}$

# Byte Ordering Functions

```
uint32_t some_long = 10;
uint16_t some_short = 20;

uint32_t network_byte_order;

// convert and send
network_byte_order = htonl(some_long);
send(s, &network_byte_order, sizeof(uint32_t), 0);

some_short == ntohs(htons(some_short)); // this expression is true
```

# inet\_aton and inet\_ntoa Functions

- These functions convert Internet addresses between

- ASCII strings - what humans prefer to use
- Network byte ordered binary values - values that are stored in socket address structures

- `inet_aton`:

- `int inet_aton(const char *strptr, struct in_addr *addrptr);`
- converting from a dots-and-numbers string into a network address in a `struct in_addr`

- `inet_ntoa`:

- `char *inet_ntoa(struct in_addr inaddr);`
- converts a 32-bit binary network byte ordered IPv4 address into its corresponding dotted-decimal (dots-and-numbers format) string.

# inet\_aton and inet\_ntoa Functions

```
struct sockaddr_in antelope;  
char *some_addr;  
  
inet_aton("10.0.0.1", &antelope.sin_addr); // store IP in antelope  
  
some_addr = inet_ntoa(antelope.sin_addr); // return the IP  
printf("%s\n", some_addr); // prints "10.0.0.1"  
  
// and this call is the same as the inet_aton() call, above:  
antelope.sin_addr.s_addr = inet_addr("10.0.0.1");
```

# inet\_pton and inet\_ntop Functions

- These two functions are new with IPv6 and work with both IPv4 and IPv6 addresses
- The letters "p" and "n" stand for presentation and numeric
  - The presentation format for an address is often an ASCII string
  - the numeric format is the binary value that goes into a socket address structure
- `int inet_pton(int family, const char *strptr, void *addrptr);`
- `const char *inet_ntop(int family, const void *addrptr, char *strptr, size_t len);`

```
#define INET_ADDRSTRLEN    16    /* for IPv4 dotted-decimal */  
#define INET6_ADDRSTRLEN  46    /* for IPv6 hex string */
```



# inet\_pton and inet\_ntop Functions

```
// IPv4 demo of inet_ntop() and inet_pton()

struct sockaddr_in sa;
char str[INET_ADDRSTRLEN];

// store this IP address in sa:
inet_pton(AF_INET, "192.0.2.33", &(sa.sin_addr));

// now get it back and print it
inet_ntop(AF_INET, &(sa.sin_addr), str, INET_ADDRSTRLEN);

printf("%s\n", str); // prints "192.0.2.33"
```

```
// IPv6 demo of inet_ntop() and inet_pton()
// (basically the same except with a bunch of 6s thrown around)

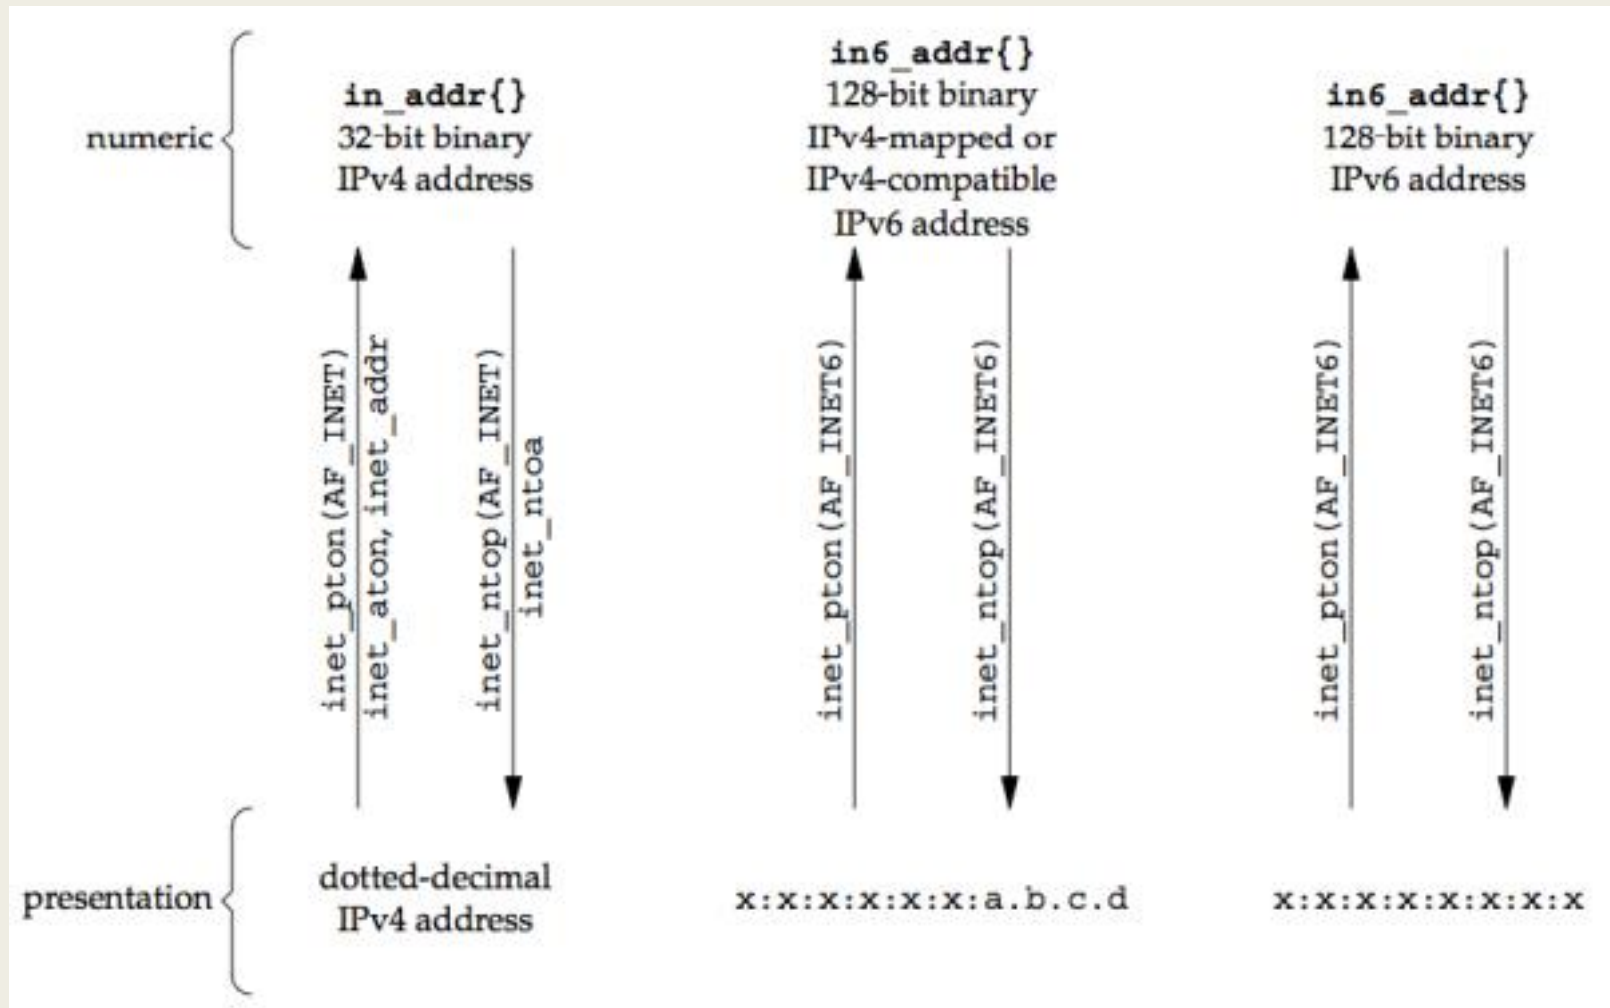
struct sockaddr_in6 sa;
char str[INET6_ADDRSTRLEN];

// store this IP address in sa:
inet_pton(AF_INET6, "2001:db8:8714:3a90::12", &(sa.sin6_addr));

// now get it back and print it
inet_ntop(AF_INET6, &(sa.sin6_addr), str, INET6_ADDRSTRLEN);

printf("%s\n", str); // prints "2001:db8:8714:3a90::12"
```

# 4 functions on address conversion



# byte\_order.c

```
printf("My IP: 111.112.221.222\n");
inet_pton(AF_INET, "111.112.221.222", &(sin.sin_addr));
printf("Original: %i\n", sin.sin_addr);

printf("Using inet_htonl: %i\n%i\n", htonl(inet_addr("111.112.221.222")),
      htonl(sin.sin_addr.s_addr));
printf("Using inet_ntohl: %i\n%i\n", ntohl(inet_addr("111.112.221.222")),
      ntohl(sin.sin_addr.s_addr));
printf("They are the same functions????\n\n");
```

```
My IP: 111.112.221.222
Original: 0xdedd706f
Using inet_htonl: 0x6f70ddde
0x6f70ddde
Using inet_ntohl: 0x6f70ddde
0x6f70ddde
They are the same functions????

Using inet_ntop: Client IP 111.112.221.222
Using inet_ntoa: Client IP 111.112.221.222
```

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
#include <netinet/in.h>
```

```
#include <arpa/inet.h>
```

```
#include <netdb.h>
```

```
#include <unistd.h>
```

```
#include <errno.h>
```

```
#include <stdio.h>
```

```
#include <string.h>
```

```
int main() {
```

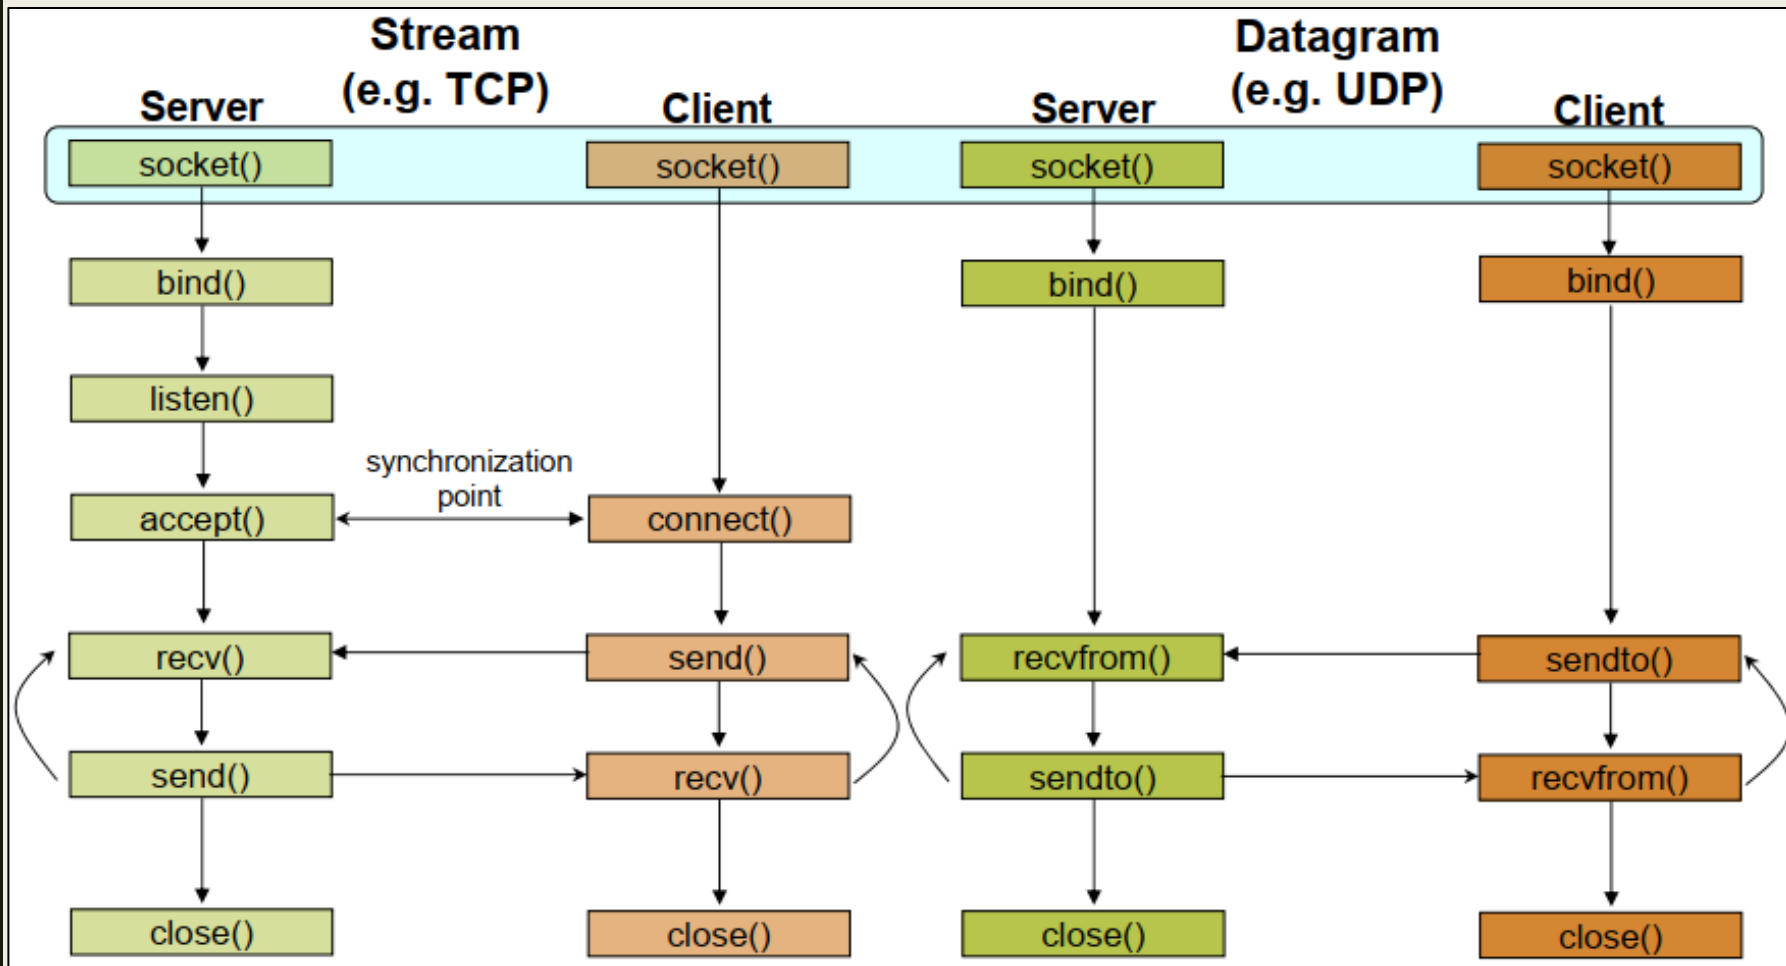
```
    struct sockaddr_in sin;
```

```
    struct sockaddr_storage client_address;
```

```
    struct sockaddr_in *sin_storage;
```

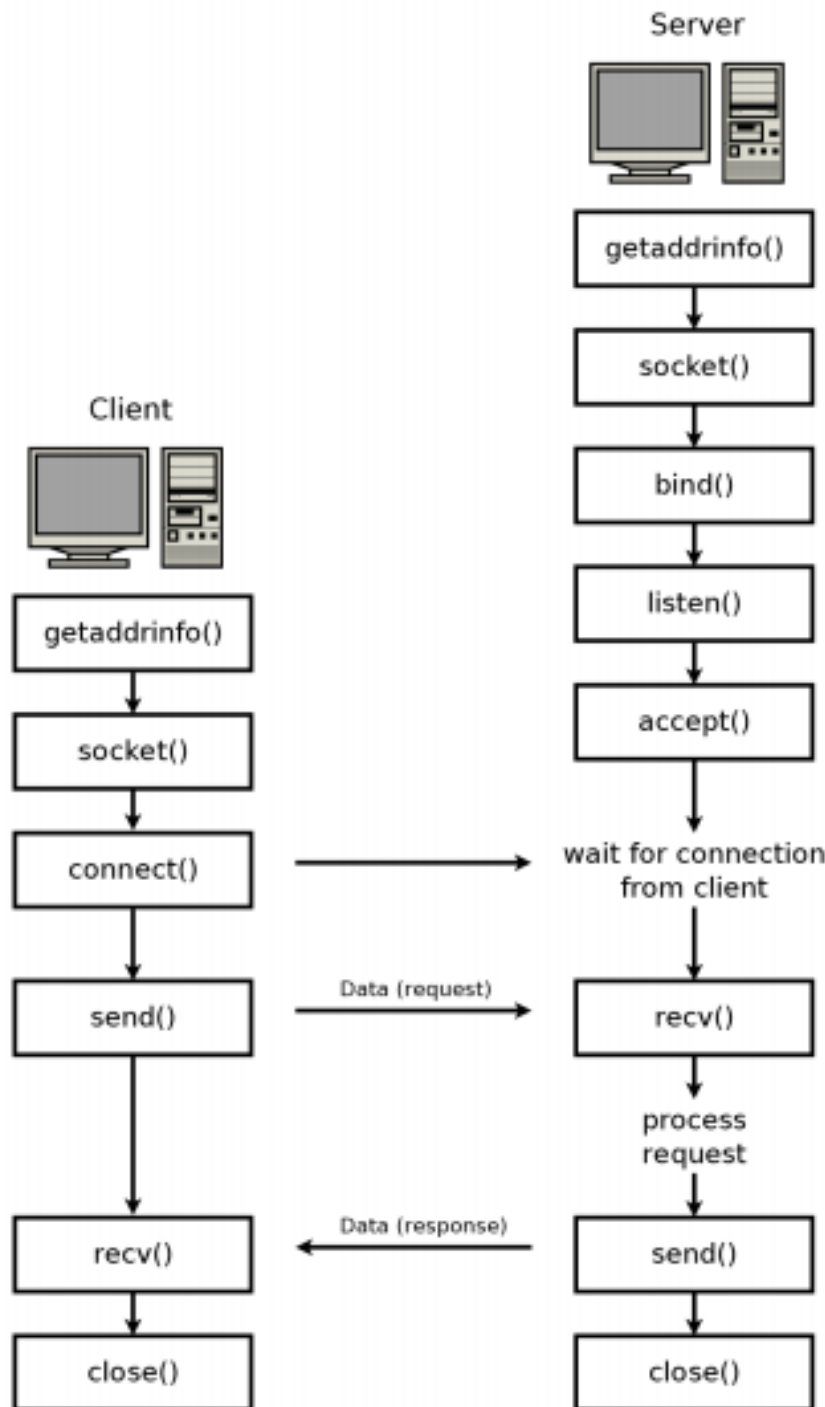
```
    char ipv4[INET_ADDRSTRLEN]
```

# Client - Server Communication



# TCP: client & server

1. The server then creates the socket with a call to `socket()`
  2. The socket must be bound to the listening IP address and port - call to `bind()`
  3. then calls `listen()`, which puts the socket in a state where it listens for new connections
  4. then call `accept()`, which will wait until a client establishes a connection to the server
  5. When the new connection has been established, `accept()` returns a *new* socket
  6. This *new* socket can be used to exchange data with the client using `send()` and `recv()`
- \*\*Meanwhile, the first socket remains listening for new connections, and repeated calls to `accept()`**



# socket()

- `int sockid = socket(family, type, protocol);`
  - **sockid**: socket descriptor, an integer (like a file-handle)
  - **family**: integer, communication domain, e.g.,
    - PF\_INET, IPv4 protocols, Internet addresses (typically used)
    - PF\_UNIX, Local communication, File addresses
  - **type**: communication type
    - SOCK\_STREAM - reliable, 2-way, connection-based service
    - SOCK\_DGRAM - unreliable, connectionless, messages of maximum length
  - **protocol**: specifies protocol
    - IPPROTO\_TCP IPPROTO\_UDP
    - usually set to 0 (i.e., use default protocol)
  - upon failure returns -1
- ☞ NOTE: socket call does not specify where data will be coming from, nor where it will be going to – it just creates the interface!

# close()

- When finished using a socket, the socket should be closed
- `status = close(sockid);`
  - **sockid**: the file descriptor (socket being closed)
  - **status**: 0 if successful, -1 if error
- Closing a socket
  - closes a connection (for stream socket)
  - frees up the port used by the socket



# bind()

- associates and reserves a port for use by the socket

- `int status = bind(sockid, &addrport, size);`
  - **sockid**: integer, socket descriptor
  - **addrport**: struct sockaddr, the (IP) address and port of the machine
    - for TCP/IP server, internet address is usually set to INADDR\_ANY, i.e., chooses any incoming interface
  - **size**: the size (in bytes) of the addrport structure
  - **status**: upon failure -1 is returned

```
int sockid;
struct sockaddr_in addrport;
sockid = socket(PF_INET, SOCK_STREAM, 0);

addrport.sin_family = AF_INET;
addrport.sin_port = htons(5100);
addrport.sin_addr.s_addr = htonl(INADDR_ANY);
if(bind(sockid, (struct sockaddr *) &addrport, sizeof(addrport)) != -1) {
    ...
}
```

# Skipping the bind()

- bind can be skipped for both types of sockets

- Datagram socket:

- if only sending, no need to bind. The OS finds a port each time the socket sends a packet
- if receiving, need to bind

- Stream socket:

- destination determined during connection setup
- don't need to know port sending from (during connection setup, receiving end is informed of port)

# listen for connections: listen()

- Instructs TCP protocol implementation to listen for connections

- `int status = listen(sockid, queueLimit);`
  - **sockid**: integer, socket descriptor
  - **queueLen**: integer, # of active participants that can “wait” for a connection
  - **status**: 0 if listening, -1 if error
- `listen()` is **non-blocking**: returns immediately
- The listening socket (sockid)
  - is never used for sending and receiving
  - is used by the server only as a way to get new sockets

# Establish Connection: Connect()

- The client establishes a connection with the server by calling connect()

- `int status = connect(sockid, &foreignAddr, addrlen);`
  - `sockid`: integer, socket to be used in connection
  - `foreignAddr`: struct sockaddr: address of the passive participant
  - `addrlen`: integer, sizeof(name)
  - status: 0 if successful connect, -1 otherwise
- `connect ()` is **blocking**

# Incoming Connection: accept()

- The server gets a socket for an incoming client connection by calling accept()

- `int s = accept(sockid, &clientAddr, &addrLen);`

- `s`: integer, the new socket (used for data-transfer)
- `sockid`: integer, the orig. socket (being listened on)
- `clientAddr`: struct sockaddr, address of the active participant
  - filled in upon return
- `addrLen`: sizeof(clientAddr): value/result parameter
  - must be set appropriately before call
  - adjusted upon return

- `accept()`

- is **blocking**: waits for connection before returning
- dequeues the next connection on the queue for socket (sockid)

# Exchanging data with stream socket

■ `int count = send(sockid, msg, msgLen, flags);`

- `msg`: `const void[]`, message to be transmitted
- `msgLen`: integer, length of message (in bytes) to transmit
- `flags`: integer, special options, usually just 0
- `count`: # bytes transmitted (-1 if error)

■ `int count = recv(sockid, recvBuf, bufLen, flags);`

- `recvBuf`: `void[]`, stores received bytes
- `bufLen`: # bytes received
- `flags`: integer, special options, usually just 0
- `count`: # bytes received (-1 if error)

■ Calls are blocking

- returns only after data is sent / received

# Reference

- Introduction to Sockets Programming in C using TCP/IP
  - Panagiota Fatourou